

# top\_25\_pandas\_tricks

2023 年 4 月 2 日

## 1 Data School: My top 25 pandas tricks ([video](#))

See also: “21 more pandas tricks” [video](#) and [notebook](#)

- Watch the [complete pandas video series](#)
- Connect on [Twitter](#), [Facebook](#), and [LinkedIn](#)
- Subscribe on [YouTube](#)
- Join the [email newsletter](#)

### 1.1 Table of contents

1. Show installed versions
2. Create an example DataFrame
3. Rename columns
4. Reverse row order
5. Reverse column order
6. Select columns by data type
7. Convert strings to numbers
8. Reduce DataFrame size
9. Build a DataFrame from multiple files (row-wise)
10. Build a DataFrame from multiple files (column-wise)
11. Create a DataFrame from the clipboard
12. Split a DataFrame into two random subsets
13. Filter a DataFrame by multiple categories
14. Filter a DataFrame by largest categories
15. Handle missing values
16. Split a string into multiple columns
17. Expand a Series of lists into a DataFrame
18. Aggregate by multiple functions
19. Combine the output of an aggregation with a DataFrame
20. Select a slice of rows and columns

21. Reshape a MultiIndexed Series
22. Create a pivot table
23. Convert continuous data into categorical data
24. Change display options
25. Style a DataFrame
26. Bonus trick: Profile a DataFrame

## 1.2 Load example datasets

```
[ ]: import pandas as pd
import numpy as np

[ ]: drinks = pd.read_csv('http://bit.ly/drinksbycountry')
movies = pd.read_csv('http://bit.ly/imdbratings')
orders = pd.read_csv('http://bit.ly/chiporders', sep='\t')
orders['item_price'] = orders.item_price.str.replace('$', '').astype('float')
stocks = pd.read_csv('http://bit.ly/smallstocks', parse_dates=['Date'])
titanic = pd.read_csv('http://bit.ly/kaggletrain')
ufo = pd.read_csv('http://bit.ly/uforeports', parse_dates=['Time'])
```

/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/ipykernel\_launcher.py:4: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expressions will *not* be treated as literal strings when regex=True.  
after removing the cwd from sys.path.

## 1.3 1. Show installed versions

Sometimes you need to know the pandas version you're using, especially when reading the pandas documentation. You can show the pandas version by typing:

```
[ ]: pd.__version__
```

```
[ ]: '1.3.5'
```

But if you also need to know the versions of pandas' dependencies, you can use the `show_versions()` function:

```
[ ]: ? pd.show_versions()
```

Object `pd.show\_versions()` not found.

You can see the versions of Python, pandas, NumPy, matplotlib, and more.

## 1.4 2. Create an example DataFrame

Let's say that you want to demonstrate some pandas code. You need an example DataFrame to work with.

There are many ways to do this, but my favorite way is to pass a dictionary to the DataFrame constructor, in which the dictionary keys are the column names and the dictionary values are lists of column values:

```
[ ]: df = pd.DataFrame({'col one': [100, 200], 'col two': [300, 400]})
df
```

```
[ ]:      col one  col two
0         100      300
1         200      400
```

Now if you need a much larger DataFrame, the above method will require way too much typing. In that case, you can use NumPy's `random.rand()` function, tell it the number of rows and columns, and pass that to the DataFrame constructor:

```
[ ]: pd.DataFrame(np.random.rand(4, 8))
```

```
[ ]:      0         1         2         3         4         5         6 \
0  0.320528  0.011367  0.583528  0.561701  0.371614  0.577766  0.705726
1  0.617830  0.212650  0.950570  0.484499  0.200131  0.887658  0.317573
2  0.680907  0.131555  0.998704  0.123795  0.743099  0.592608  0.898709
3  0.284402  0.369183  0.587917  0.634587  0.125579  0.349900  0.638982

      7
0  0.841400
1  0.587168
2  0.138611
3  0.791359
```

That's pretty good, but if you also want non-numeric column names, you can coerce a string of letters to a list and then pass that list to the `columns` parameter:

```
[ ]: pd.DataFrame(np.random.rand(4, 8), columns=list('abcdefgh'))
```

```
[ ]:      a         b         c         d         e         f         g \
0  0.338772  0.636219  0.364416  0.791293  0.680778  0.649622  0.200556
1  0.917255  0.875851  0.898739  0.072976  0.186387  0.748383  0.203719
2  0.735625  0.296076  0.465671  0.161475  0.255853  0.785485  0.949005
3  0.854833  0.406767  0.184608  0.719145  0.810886  0.661914  0.462479
```

```
      h
0  0.587630
1  0.442413
2  0.985340
3  0.309318
```

As you might guess, your string will need to have the same number of characters as there are columns.

### 1.5 3. Rename columns

Let's take a look at the example DataFrame we created in the last trick:

```
[ ]: df
      col one  col two
0         100      300
1         200      400
```

I prefer to use dot notation to select pandas columns, but that won't work since the column names have spaces. Let's fix this.

The most flexible method for renaming columns is the `rename()` method. You pass it a dictionary in which the keys are the old names and the values are the new names, and you also specify the axis:

```
[ ]: df = df.rename({'col one': 'col_one', 'col two': 'col_two'}, axis='columns')
```

The best thing about this method is that you can use it to rename any number of columns, whether it be just one column or all columns.

Now if you're going to rename all of the columns at once, a simpler method is just to overwrite the columns attribute of the DataFrame:

```
[ ]: df.columns = ['col_one', 'col_two']
```

Now if the only thing you're doing is replacing spaces with underscores, an even better method is to use the `str.replace()` method, since you don't have to type out all of the column names:

```
[ ]: df.columns = df.columns.str.replace(' ', '_')
```

All three of these methods have the same result, which is to rename the columns so that they don't have any spaces:

```
[ ]: df
```

```
[ ]:   col_one  col_two
0      100     300
1      200     400
```

Finally, if you just need to add a prefix or suffix to all of your column names, you can use the `add_prefix()` method...

```
[ ]: df.add_prefix('X_')
```

```
[ ]:   X_col_one  X_col_two
0         100         300
1         200         400
```

...or the `add_suffix()` method:

```
[ ]: df.add_suffix('_Y')
```

```
[ ]:   col_one_Y  col_two_Y
0         100         300
1         200         400
```

## 1.6 4. Reverse row order

Let's take a look at the drinks DataFrame:

```
[ ]: drinks.head()
```

```
[ ]:   country  beer_servings  spirit_servings  wine_servings  \
0  Afghanistan           0              0              0
1    Albania           89             132             54
2    Algeria           25              0             14
3   Andorra          245             138            312
4    Angola          217              57             45
```

```
   total_litres_of_pure_alcohol  continent
0              0.0      Asia
1              4.9    Europe
2              0.7    Africa
3             12.4    Europe
4              5.9    Africa
```

This is a dataset of average alcohol consumption by country. What if you wanted to reverse the order of the rows?

The most straightforward method is to use the `loc` accessor and pass it `::-1`, which is the same slicing notation used to reverse a Python list:

```
[ ]: drinks.loc[::-1].head()
```

```
[ ]:      country  beer_servings  spirit_servings  wine_servings  \
192  Zimbabwe          64           18           4
191   Zambia          32           19           4
190    Yemen           6            0           0
189  Vietnam         111            2           1
188  Venezuela        333          100           3

      total_litres_of_pure_alcohol  continent
192                          4.7      Africa
191                          2.5      Africa
190                          0.1       Asia
189                          2.0       Asia
188                          7.7  South America
```

What if you also wanted to reset the index so that it starts at zero?

You would use the `reset_index()` method and tell it to drop the old index entirely:

```
[ ]: drinks.loc[::-1].reset_index(drop=True).head()
```

```
[ ]:      country  beer_servings  spirit_servings  wine_servings  \
0  Zimbabwe          64           18           4
1   Zambia          32           19           4
2    Yemen           6            0           0
3  Vietnam         111            2           1
4  Venezuela        333          100           3

      total_litres_of_pure_alcohol  continent
0                          4.7      Africa
1                          2.5      Africa
2                          0.1       Asia
3                          2.0       Asia
4                          7.7  South America
```

As you can see, the rows are in reverse order but the index has been reset to the default integer

index.

## 1.7 5. Reverse column order

Similar to the previous trick, you can also use `loc` to reverse the left-to-right order of your columns:

```
[ ]: drinks.loc[:, ::-1].head()
```

```
[ ]:  continent  total_litres_of_pure_alcohol  wine_servings  spirit_servings  \
0      Asia                0.0                0                0
1    Europe                4.9               54             132
2    Africa                0.7               14                0
3    Europe               12.4              312             138
4    Africa                5.9               45             57

      beer_servings  country
0                0  Afghanistan
1               89    Albania
2               25    Algeria
3              245    Andorra
4              217    Angola
```

The colon before the comma means “select all rows” , and the `::-1` after the comma means “reverse the columns” , which is why “country” is now on the right side.

## 1.8 6. Select columns by data type

Here are the data types of the drinks DataFrame:

```
[ ]: drinks.dtypes
```

```
[ ]: country                object
      beer_servings         int64
      spirit_servings        int64
      wine_servings         int64
      total_litres_of_pure_alcohol  float64
      continent              object
      dtype: object
```

Let’ s say you need to select only the numeric columns. You can use the `select_dtypes()` method:

```
[ ]: drinks.select_dtypes(include='number').head()
```

```
[ ]:   beer_servings  spirit_servings  wine_servings  total_litres_of_pure_alcohol
0           0           0           0           0.0
1          89          132           54           4.9
2          25           0           14           0.7
3         245          138          312          12.4
4         217           57           45           5.9
```

This includes both int and float columns.

You could also use this method to select just the object columns:

```
[ ]: drinks.select_dtypes(include='object').head()
```

```
[ ]:   country continent
0  Afghanistan    Asia
1    Albania    Europe
2    Algeria    Africa
3   Andorra    Europe
4    Angola    Africa
```

You can tell it to include multiple data types by passing a list:

```
[ ]: drinks.select_dtypes(include=['number', 'object', 'category', 'datetime']).head()
```

```
[ ]:   country  beer_servings  spirit_servings  wine_servings  \
0  Afghanistan           0           0           0
1    Albania           89          132           54
2    Algeria           25           0           14
3   Andorra          245          138          312
4    Angola           217           57           45

   total_litres_of_pure_alcohol  continent
0                0.0        Asia
1                4.9        Europe
2                0.7        Africa
3               12.4        Europe
4                5.9        Africa
```

You can also tell it to exclude certain data types:

```
[ ]: drinks.select_dtypes(exclude='number').head()
```



```
[ ]:      country continent
0  Afghanistan      Asia
1     Albania    Europe
2     Algeria    Africa
3     Andorra    Europe
4      Angola    Africa
```

## 1.9 7. Convert strings to numbers

Let's create another example DataFrame:

```
[ ]: df = pd.DataFrame({'col_one': ['1.1', '2.2', '3.3'],
                        'col_two': ['4.4', '5.5', '6.6'],
                        'col_three': ['7.7', '8.8', '-']})
df
```

```
[ ]:   col_one col_two col_three
0      1.1     4.4      7.7
1      2.2     5.5      8.8
2      3.3     6.6      -
```

These numbers are actually stored as strings, which results in object columns:

```
[ ]: df.dtypes
```

```
[ ]: col_one      object
     col_two      object
     col_three      object
     dtype: object
```

In order to do mathematical operations on these columns, we need to convert the data types to numeric. You can use the `astype()` method on the first two columns:

```
[ ]: df.astype({'col_one': 'float', 'col_two': 'float'}).dtypes
```

```
[ ]: col_one      float64
     col_two      float64
     col_three      object
     dtype: object
```

However, this would have resulted in an error if you tried to use it on the third column, because that column contains a dash to represent zero and pandas doesn't understand how to handle it.

Instead, you can use the `to_numeric()` function on the third column and tell it to convert any invalid input into `NaN` values:

```
[ ]: pd.to_numeric(df.col_three, errors='coerce')
```

```
[ ]: 0    7.7
      1    8.8
      2   NaN
      Name: col_three, dtype: float64
```

If you know that the `NaN` values actually represent zeros, you can fill them with zeros using the `fillna()` method:

```
[ ]: pd.to_numeric(df.col_three, errors='coerce').fillna(0)
```

```
[ ]: 0    7.7
      1    8.8
      2    0.0
      Name: col_three, dtype: float64
```

Finally, you can apply this function to the entire DataFrame all at once by using the `apply()` method:

```
[ ]: df = df.apply(pd.to_numeric, errors='coerce').fillna(0)
      df
```

```
[ ]:   col_one  col_two  col_three
      0     1.1     4.4         7.7
      1     2.2     5.5         8.8
      2     3.3     6.6         0.0
```

This one line of code accomplishes our goal, because all of the data types have now been converted to float:

```
[ ]: df.dtypes
```

```
[ ]: col_one      float64
      col_two      float64
      col_three      float64
      dtype: object
```

## 1.10 8. Reduce DataFrame size

pandas DataFrames are designed to fit into memory, and so sometimes you need to reduce the DataFrame size in order to work with it on your system.

Here's the size of the drinks DataFrame:

```
[ ]: drinks.info(memory_usage='deep')
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 193 entries, 0 to 192
Data columns (total 6 columns):
country                193 non-null object
beer_servings          193 non-null int64
spirit_servings        193 non-null int64
wine_servings          193 non-null int64
total_litres_of_pure_alcohol  193 non-null float64
continent              193 non-null object
dtypes: float64(1), int64(3), object(2)
memory usage: 30.4 KB
```

You can see that it currently uses 30.4 KB.

If you're having performance problems with your DataFrame, or you can't even read it into memory, there are two easy steps you can take during the file reading process to reduce the DataFrame size.

The first step is to only read in the columns that you actually need, which we specify with the "usecols" parameter:

```
[ ]: cols = ['beer_servings', 'continent']
small_drinks = pd.read_csv('http://bit.ly/drinksbycountry', usecols=cols)
small_drinks.info(memory_usage='deep')
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 193 entries, 0 to 192
Data columns (total 2 columns):
beer_servings    193 non-null int64
continent        193 non-null object
dtypes: int64(1), object(1)
memory usage: 13.6 KB
```

By only reading in these two columns, we've reduced the DataFrame size to 13.6 KB.

The second step is to convert any object columns containing categorical data to the category data type, which we specify with the "dtype" parameter:

```
[ ]: dtypes = {'continent': 'category'}
smaller_drinks = pd.read_csv('http://bit.ly/drinksbycountry', usecols=cols,
    ↪ dtype=dtypes)
```

```
smaller_drinks.info(memory_usage='deep')
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 193 entries, 0 to 192
```

```
Data columns (total 2 columns):
```

```
beer_servings    193 non-null int64
```

```
continent        193 non-null category
```

```
dtypes: category(1), int64(1)
```

```
memory usage: 2.3 KB
```

By reading in the continent column as the category data type, we've further reduced the DataFrame size to 2.3 KB.

Keep in mind that the category data type will only reduce memory usage if you have a small number of categories relative to the number of rows.

## 1.11 9. Build a DataFrame from multiple files (row-wise)

Let's say that your dataset is spread across multiple files, but you want to read the dataset into a single DataFrame.

For example, I have a small dataset of stock data in which each CSV file only includes a single day. Here's the first day:

```
[ ]: pd.read_csv('data/stocks1.csv')
```

```
[ ]:      Date    Close    Volume Symbol
0  2016-10-03    31.50   14070500   CSC0
1  2016-10-03   112.52   21701800  AAPL
2  2016-10-03    57.42   19189500  MSFT
```

Here's the second day:

```
[ ]: pd.read_csv('data/stocks2.csv')
```

```
[ ]:      Date    Close    Volume Symbol
0  2016-10-04   113.00   29736800  AAPL
1  2016-10-04    57.24   20085900  MSFT
2  2016-10-04    31.35   18460400   CSC0
```

And here's the third day:

```
[ ]: pd.read_csv('data/stocks3.csv')
```

```
[ ]:      Date    Close    Volume Symbol
0  2016-10-05    57.64  16726400   MSFT
1  2016-10-05    31.59  11808600   CSCD
2  2016-10-05   113.05  21453100   AAPL
```

You could read each CSV file into its own DataFrame, combine them together, and then delete the original DataFrames, but that would be memory inefficient and require a lot of code.

A better solution is to use the built-in glob module:

```
[ ]: from glob import glob
```

You can pass a pattern to `glob()`, including wildcard characters, and it will return a list of all files that match that pattern.

In this case, glob is looking in the “data” subdirectory for all CSV files that start with the word “stocks” :

```
[ ]: stock_files = sorted(glob('data/stocks*.csv'))
stock_files
```

```
[ ]: ['data/stocks1.csv', 'data/stocks2.csv', 'data/stocks3.csv']
```

glob returns filenames in an arbitrary order, which is why we sorted the list using Python’s built-in `sorted()` function.

We can then use a generator expression to read each of the files using `read_csv()` and pass the results to the `concat()` function, which will concatenate the rows into a single DataFrame:

```
[ ]: pd.concat((pd.read_csv(file) for file in stock_files))
```

```
[ ]:      Date    Close    Volume Symbol
0  2016-10-03    31.50  14070500   CSCD
1  2016-10-03   112.52  21701800   AAPL
2  2016-10-03    57.42  19189500   MSFT
0  2016-10-04   113.00  29736800   AAPL
1  2016-10-04    57.24  20085900   MSFT
2  2016-10-04    31.35  18460400   CSCD
0  2016-10-05    57.64  16726400   MSFT
1  2016-10-05    31.59  11808600   CSCD
2  2016-10-05   113.05  21453100   AAPL
```

Unfortunately, there are now duplicate values in the index. To avoid that, we can tell the `concat()` function to ignore the index and instead use the default integer index:

```
[ ]: pd.concat((pd.read_csv(file) for file in stock_files), ignore_index=True)
```

```
[ ]:
      Date    Close    Volume Symbol
0  2016-10-03   31.50  14070500   CSCO
1  2016-10-03  112.52  21701800   AAPL
2  2016-10-03   57.42  19189500   MSFT
3  2016-10-04  113.00  29736800   AAPL
4  2016-10-04   57.24  20085900   MSFT
5  2016-10-04   31.35  18460400   CSCO
6  2016-10-05   57.64  16726400   MSFT
7  2016-10-05   31.59  11808600   CSCO
8  2016-10-05  113.05  21453100   AAPL
```

## 1.12 10. Build a DataFrame from multiple files (column-wise)

The previous trick is useful when each file contains rows from your dataset. But what if each file instead contains columns from your dataset?

Here's an example in which the drinks dataset has been split into two CSV files, and each file contains three columns:

```
[ ]: pd.read_csv('data/drinks1.csv').head()
```

```
[ ]:
      country  beer_servings  spirit_servings
0  Afghanistan           0             0
1    Albania           89           132
2    Algeria           25             0
3   Andorra          245           138
4    Angola          217           57
```

```
[ ]: pd.read_csv('data/drinks2.csv').head()
```

```
[ ]:
      wine_servings  total_litres_of_pure_alcohol  continent
0             0                                0.0      Asia
1             54                                4.9    Europe
2             14                                0.7    Africa
3            312                               12.4    Europe
4             45                                5.9    Africa
```

Similar to the previous trick, we'll start by using `glob()`:

```
[ ]: drink_files = sorted(glob('data/drinks*.csv'))
```

And this time, we'll tell the `concat()` function to concatenate along the columns axis:

```
[ ]: pd.concat((pd.read_csv(file) for file in drink_files), axis='columns').head()
```

```
[ ]:      country  beer_servings  spirit_servings  wine_servings  \
0  Afghanistan          0           0           0
1    Albania          89          132          54
2    Algeria          25           0          14
3    Andorra         245          138         312
4     Angola         217           57          45

      total_litres_of_pure_alcohol  continent
0                0.0        Asia
1                4.9        Europe
2                0.7        Africa
3               12.4        Europe
4                5.9        Africa
```

Now our DataFrame has all six columns.

### 1.13 11. Create a DataFrame from the clipboard

Let's say that you have some data stored in an Excel spreadsheet or a [Google Sheet](#), and you want to get it into a DataFrame as quickly as possible.

Just select the data and copy it to the clipboard. Then, you can use the `read_clipboard()` function to read it into a DataFrame:

```
[ ]: df = pd.read_clipboard()
df
```

```
[ ]:      Column A  Column B  Column C
0         1        4.4    seven
1         2        5.5    eight
2         3        6.6     nine
```

Just like the `read_csv()` function, `read_clipboard()` automatically detects the correct data type for each column:

```
[ ]: df.dtypes
```

```
[ ]: Column A      int64
      Column B    float64
      Column C     object
```

dtype: object

Let's copy one other dataset to the clipboard:

```
[ ]: df = pd.read_clipboard()
df
```

```
[ ]:      Left  Right
Alice     10     40
Bob       20     50
Charlie   30     60
```

Amazingly, pandas has even identified the first column as the index:

```
[ ]: df.index
```

```
[ ]: Index(['Alice', 'Bob', 'Charlie'], dtype='object')
```

Keep in mind that if you want your work to be reproducible in the future, `read_clipboard()` is not the recommended approach.

## 1.14 12. Split a DataFrame into two random subsets

Let's say that you want to split a DataFrame into two parts, randomly assigning 75% of the rows to one DataFrame and the other 25% to a second DataFrame.

For example, we have a DataFrame of movie ratings with 979 rows:

```
[ ]: len(movies)
```

```
[ ]: 979
```

We can use the `sample()` method to randomly select 75% of the rows and assign them to the “movies\_1” DataFrame:

```
[ ]: movies_1 = movies.sample(frac=0.75, random_state=1234)
```

Then we can use the `drop()` method to drop all rows that are in “movies\_1” and assign the remaining rows to “movies\_2” :

```
[ ]: movies_2 = movies.drop(movies_1.index)
```

You can see that the total number of rows is correct:

```
[ ]: len(movies_1) + len(movies_2)
```

```
[ ]: 979
```



And you can see from the index that every movie is in either “movies\_1” :

```
[ ]: movies_1.index.sort_values()
```

```
[ ]: Int64Index([ 0,  2,  5,  6,  7,  8,  9, 11, 13, 16,
               ...
               966, 967, 969, 971, 972, 974, 975, 976, 977, 978],
               dtype='int64', length=734)
```

...or “movies\_2” :

```
[ ]: movies_2.index.sort_values()
```

```
[ ]: Int64Index([ 1,  3,  4, 10, 12, 14, 15, 18, 26, 30,
               ...
               931, 934, 937, 941, 950, 954, 960, 968, 970, 973],
               dtype='int64', length=245)
```

Keep in mind that this approach will not work if your index values are not unique.

## 1.15 13. Filter a DataFrame by multiple categories

Let’ s take a look at the movies DataFrame:

```
[ ]: movies.head()
```

```
[ ]:   star_rating      title content_rating  genre  duration \
0      9.3  The Shawshank Redemption          R   Crime      142
1      9.2      The Godfather              R   Crime      175
2      9.1  The Godfather: Part II          R   Crime      200
3      9.0      The Dark Knight          PG-13  Action      152
4      8.9      Pulp Fiction              R   Crime      154
```

```
                                actors_list
0  [u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt...
1  [u'Marlon Brando', u'Al Pacino', u'James Caan']
2  [u'Al Pacino', u'Robert De Niro', u'Robert Duv...
3  [u'Christian Bale', u'Heath Ledger', u'Aaron E...
4  [u'John Travolta', u'Uma Thurman', u'Samuel L...
```

One of the columns is genre:

```
[ ]: movies.genre.unique()
```

```
[ ]: array(['Crime', 'Action', 'Drama', 'Western', 'Adventure', 'Biography',
          'Comedy', 'Animation', 'Mystery', 'Horror', 'Film-Noir', 'Sci-Fi',
          'History', 'Thriller', 'Family', 'Fantasy'], dtype=object)
```

If we wanted to filter the DataFrame to only show movies with the genre Action or Drama or Western, we could use multiple conditions separated by the “or” operator:

```
[ ]: movies[(movies.genre == 'Action') |
            (movies.genre == 'Drama') |
            (movies.genre == 'Western')].head()
```

```
[ ]:      star_rating      title content_rating  genre \
3         9.0      The Dark Knight      PG-13  Action
5         8.9      12 Angry Men    NOT RATED   Drama
6         8.9  The Good, the Bad and the Ugly    NOT RATED  Western
9         8.9      Fight Club          R     Drama
11        8.8      Inception      PG-13    Action

      duration      actors_list
3         152  [u'Christian Bale', u'Heath Ledger', u'Aaron E...
5          96  [u'Henry Fonda', u'Lee J. Cobb', u'Martin Bals...
6         161  [u'Clint Eastwood', u'Eli Wallach', u'Lee Van ...
9         139  [u'Brad Pitt', u'Edward Norton', u'Helena Bonh...
11        148  [u'Leonardo DiCaprio', u'Joseph Gordon-Levitt'...
```

However, you can actually rewrite this code more clearly by using the `isin()` method and passing it a list of genres:

```
[ ]: movies[movies.genre.isin(['Action', 'Drama', 'Western'])].head()
```

```
[ ]:      star_rating      title content_rating  genre \
3         9.0      The Dark Knight      PG-13  Action
5         8.9      12 Angry Men    NOT RATED   Drama
6         8.9  The Good, the Bad and the Ugly    NOT RATED  Western
9         8.9      Fight Club          R     Drama
11        8.8      Inception      PG-13    Action

      duration      actors_list
3         152  [u'Christian Bale', u'Heath Ledger', u'Aaron E...
5          96  [u'Henry Fonda', u'Lee J. Cobb', u'Martin Bals...
6         161  [u'Clint Eastwood', u'Eli Wallach', u'Lee Van ...
9         139  [u'Brad Pitt', u'Edward Norton', u'Helena Bonh...
```

```
11          148 [u'Leonardo DiCaprio', u'Joseph Gordon-Levitt'...
```

And if you want to reverse this filter, so that you are excluding (rather than including) those three genres, you can put a tilde in front of the condition:

```
[ ]: movies[~movies.genre.isin(['Action', 'Drama', 'Western'])].head()
```

```
[ ]:      star_rating      title content_rating \
0          9.3      The Shawshank Redemption      R
1          9.2      The Godfather              R
2          9.1      The Godfather: Part II      R
4          8.9      Pulp Fiction                R
7          8.9  The Lord of the Rings: The Return of the King  PG-13
```

```
      genre  duration      actors_list
0    Crime    142  [u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt...
1    Crime    175   [u'Marlon Brando', u'Al Pacino', u'James Caan']
2    Crime    200  [u'Al Pacino', u'Robert De Niro', u'Robert Duv...
4    Crime    154  [u'John Travolta', u'Uma Thurman', u'Samuel L...
7  Adventure    201  [u'Elijah Wood', u'Viggo Mortensen', u'Ian McK...
```

This works because tilde is the “not” operator in Python.

## 1.16 14. Filter a DataFrame by largest categories

Let’s say that you needed to filter the movies DataFrame by genre, but only include the 3 largest genres.

We’ll start by taking the `value_counts()` of genre and saving it as a Series called counts:

```
[ ]: counts = movies.genre.value_counts()
      counts
```

```
[ ]: Drama      278
      Comedy    156
      Action    136
      Crime     124
      Biography   77
      Adventure   75
      Animation   62
      Horror      29
      Mystery     16
```

```

Western      9
Sci-Fi       5
Thriller     5
Film-Noir    3
Family       2
Fantasy      1
History      1
Name: genre, dtype: int64

```

The Series method `nlargest()` makes it easy to select the 3 largest values in this Series:

```
[ ]: counts.nlargest(3)
```

```

[ ]: Drama      278
      Comedy    156
      Action    136
      Name: genre, dtype: int64

```

And all we actually need from this Series is the index:

```
[ ]: counts.nlargest(3).index
```

```
[ ]: Index(['Drama', 'Comedy', 'Action'], dtype='object')
```

Finally, we can pass the index object to `isin()`, and it will be treated like a list of genres:

```
[ ]: movies[movies.genre.isin(counts.nlargest(3).index)].head()
```

```

[ ]:   star_rating      title \
3      9.0      The Dark Knight
5      8.9      12 Angry Men
9      8.9      Fight Club
11     8.8      Inception
12     8.8  Star Wars: Episode V - The Empire Strikes Back

```

```

      content_rating  genre  duration \
3      PG-13  Action      152
5     NOT RATED  Drama      96
9         R    Drama     139
11     PG-13  Action     148
12         PG  Action     124

```

actors\_list

```

3  [u'Christian Bale', u'Heath Ledger', u'Aaron E...
5  [u'Henry Fonda', u'Lee J. Cobb', u'Martin Bals...
9  [u'Brad Pitt', u'Edward Norton', u'Helena Bonh...
11 [u'Leonardo DiCaprio', u'Joseph Gordon-Levitt'...
12 [u'Mark Hamill', u'Harrison Ford', u'Carrie Fi...

```

Thus, only Drama and Comedy and Action movies remain in the DataFrame.

## 1.17 15. Handle missing values

Let's look at a dataset of UFO sightings:

```
[ ]: ufo.head()
```

```
[ ]:
      City Colors Reported Shape Reported State \
0      Ithaca          NaN    TRIANGLE    NY
1  Willingboro          NaN      OTHER    NJ
2      Holyoke          NaN      OVAL    CO
3      Abilene          NaN      DISK    KS
4 New York Worlds Fair          NaN    LIGHT    NY

      Time
0 1930-06-01 22:00:00
1 1930-06-30 20:00:00
2 1931-02-15 14:00:00
3 1931-06-01 13:00:00
4 1933-04-18 19:00:00

```

You'll notice that some of the values are missing.

To find out how many values are missing in each column, you can use the `isna()` method and then take the `sum()`:

```
[ ]: ufo.isna().sum()
```

```
[ ]: City          25
     Colors Reported 15359
     Shape Reported  2644
     State           0
     Time            0
     dtype: int64

```

`isna()` generated a DataFrame of True and False values, and `sum()` converted all of the True values to 1 and added them up.

Similarly, you can find out the percentage of values that are missing by taking the `mean()` of `isna()`:

```
[ ]: ufo.isna().mean()
```

```
[ ]: City                0.001371
      Colors Reported    0.842004
      Shape Reported    0.144948
      State              0.000000
      Time               0.000000
      dtype: float64
```

If you want to drop the columns that have any missing values, you can use the `dropna()` method:

```
[ ]: ufo.dropna(axis='columns').head()
```

```
[ ]:   State                Time
0    NY 1930-06-01 22:00:00
1    NJ 1930-06-30 20:00:00
2    CO 1931-02-15 14:00:00
3    KS 1931-06-01 13:00:00
4    NY 1933-04-18 19:00:00
```

Or if you want to drop columns in which more than 10% of the values are missing, you can set a threshold for `dropna()`:

```
[ ]: ufo.dropna(thresh=len(ufo)*0.9, axis='columns').head()
```

```
[ ]:   City State                Time
0    Ithaca  NY 1930-06-01 22:00:00
1  Willingboro  NJ 1930-06-30 20:00:00
2    Holyoke  CO 1931-02-15 14:00:00
3    Abilene  KS 1931-06-01 13:00:00
4 New York Worlds Fair  NY 1933-04-18 19:00:00
```

`len(ufo)` returns the total number of rows, and then we multiply that by 0.9 to tell pandas to only keep columns in which at least 90% of the values are not missing.

## 1.18 16. Split a string into multiple columns

Let's create another example DataFrame:

```
[ ]: df = pd.DataFrame({'name': ['John Arthur Doe', 'Jane Ann Smith'],  
                        'location': ['Los Angeles, CA', 'Washington, DC']})  
df
```

```
[ ]:      name      location  
0  John Arthur Doe  Los Angeles, CA  
1   Jane Ann Smith  Washington, DC
```

What if we wanted to split the “name” column into three separate columns, for first, middle, and last name? We would use the `str.split()` method and tell it to split on a space character and expand the results into a DataFrame:

```
[ ]: df.name.str.split(' ', expand=True)
```

```
[ ]:      0      1      2  
0  John  Arthur  Doe  
1  Jane    Ann  Smith
```

These three columns can actually be saved to the original DataFrame in a single assignment statement:

```
[ ]: df[['first', 'middle', 'last']] = df.name.str.split(' ', expand=True)  
df
```

```
[ ]:      name      location first  middle  last  
0  John Arthur Doe  Los Angeles, CA  John  Arthur  Doe  
1   Jane Ann Smith  Washington, DC  Jane    Ann  Smith
```

What if we wanted to split a string, but only keep one of the resulting columns? For example, let's split the location column on “comma space” :

```
[ ]: df.location.str.split(', ', expand=True)
```

```
[ ]:      0      1  
0  Los Angeles  CA  
1  Washington  DC
```

If we only cared about saving the city name in column 0, we can just select that column and save it to the DataFrame:

```
[ ]: df['city'] = df.location.str.split(', ', expand=True)[0]  
df
```

```
[ ]:      name      location first  middle  last      city
0  John Arthur Doe  Los Angeles, CA  John  Arthur   Doe  Los Angeles
1  Jane Ann Smith   Washington, DC  Jane   Ann   Smith  Washington
```

## 1.19 17. Expand a Series of lists into a DataFrame

Let's create another example DataFrame:

```
[ ]: df = pd.DataFrame({'col_one':['a', 'b', 'c'], 'col_two':[[10, 40], [20, 50], [30, 60]]})
df
```

```
[ ]:   col_one  col_two
0      a  [10, 40]
1      b  [20, 50]
2      c  [30, 60]
```

There are two columns, and the second column contains regular Python lists of integers.

If we wanted to expand the second column into its own DataFrame, we can use the `apply()` method on that column and pass it the Series constructor:

```
[ ]: df_new = df.col_two.apply(pd.Series)
df_new
```

```
[ ]:    0  1
0  10  40
1  20  50
2  30  60
```

And by using the `concat()` function, you can combine the original DataFrame with the new DataFrame:

```
[ ]: pd.concat([df, df_new], axis='columns')
```

```
[ ]:   col_one  col_two  0  1
0      a  [10, 40]  10  40
1      b  [20, 50]  20  50
2      c  [30, 60]  30  60
```

## 1.20 18. Aggregate by multiple functions

Let's look at a DataFrame of orders from the Chipotle restaurant chain:



```
[ ]: orders.head(10)
```

```
[ ]:  order_id  quantity                item_name \
0         1         1      Chips and Fresh Tomato Salsa
1         1         1                          Izze
2         1         1          Nantucket Nectar
3         1         1  Chips and Tomatillo-Green Chili Salsa
4         2         2          Chicken Bowl
5         3         1          Chicken Bowl
6         3         1          Side of Chips
7         4         1          Steak Burrito
8         4         1        Steak Soft Tacos
9         5         1          Steak Burrito

      choice_description  item_price
0                    NaN          2.39
1          [Clementine]          3.39
2            [Apple]          3.39
3                    NaN          2.39
4  [Tomatillo-Red Chili Salsa (Hot), [Black Beans...    16.98
5  [Fresh Tomato Salsa (Mild), [Rice, Cheese, Sou...    10.98
6                    NaN          1.69
7  [Tomatillo Red Chili Salsa, [Fajita Vegetables...    11.75
8  [Tomatillo Green Chili Salsa, [Pinto Beans, Ch...     9.25
9  [Fresh Tomato Salsa, [Rice, Black Beans, Pinto...     9.25
```

Each order has an `order_id` and consists of one or more rows. To figure out the total price of an order, you sum the `item_price` for that `order_id`. For example, here's the total price of order number 1:

```
[ ]: orders[orders.order_id == 1].item_price.sum()
```

```
[ ]: 11.56
```

If you wanted to calculate the total price of every order, you would `groupby()` `order_id` and then take the sum of `item_price` for each group:

```
[ ]: orders.groupby('order_id').item_price.sum().head()
```

```
[ ]: order_id
1    11.56
2    16.98
```

```

3    12.67
4    21.00
5    13.70
Name: item_price, dtype: float64

```

However, you're not actually limited to aggregating by a single function such as `sum()`. To aggregate by multiple functions, you use the `agg()` method and pass it a list of functions such as `sum()` and `count()`:

```
[ ]: orders.groupby('order_id').item_price.agg(['sum', 'count']).head()
```

```
[ ]:
      sum  count
order_id
1      11.56     4
2      16.98     1
3      12.67     2
4      21.00     2
5      13.70     2

```

That gives us the total price of each order as well as the number of items in each order.

## 1.21 19. Combine the output of an aggregation with a DataFrame

Let's take another look at the orders DataFrame:

```
[ ]: orders.head(10)
```

```
[ ]:
  order_id  quantity  item_name \
0         1         1  Chips and Fresh Tomato Salsa
1         1         1              Izze
2         1         1  Nantucket Nectar
3         1         1  Chips and Tomatillo-Green Chili Salsa
4         2         2      Chicken Bowl
5         3         1      Chicken Bowl
6         3         1      Side of Chips
7         4         1      Steak Burrito
8         4         1  Steak Soft Tacos
9         5         1      Steak Burrito

      choice_description  item_price
0                    NaN          2.39
1          [Clementine]          3.39

```

2	[Apple]	3.39
3	NaN	2.39
4	[Tomatillo-Red Chili Salsa (Hot), [Black Beans...	16.98
5	[Fresh Tomato Salsa (Mild), [Rice, Cheese, Sou...	10.98
6	NaN	1.69
7	[Tomatillo Red Chili Salsa, [Fajita Vegetables...	11.75
8	[Tomatillo Green Chili Salsa, [Pinto Beans, Ch...	9.25
9	[Fresh Tomato Salsa, [Rice, Black Beans, Pinto...	9.25

What if we wanted to create a new column listing the total price of each order? Recall that we calculated the total price using the `sum()` method:

```
[ ]: orders.groupby('order_id').item_price.sum().head()
```

```
[ ]: order_id
1    11.56
2    16.98
3    12.67
4    21.00
5    13.70
Name: item_price, dtype: float64
```

`sum()` is an aggregation function, which means that it returns a reduced version of the input data.

In other words, the output of the `sum()` function:

```
[ ]: len(orders.groupby('order_id').item_price.sum())
```

```
[ ]: 1834
```

...is smaller than the input to the function:

```
[ ]: len(orders.item_price)
```

```
[ ]: 4622
```

The solution is to use the `transform()` method, which performs the same calculation but returns output data that is the same shape as the input data:

```
[ ]: total_price = orders.groupby('order_id').item_price.transform('sum')
len(total_price)
```

```
[ ]: 4622
```

We'll store the results in a new DataFrame column called `total_price`:

```
[ ]: orders['total_price'] = total_price
orders.head(10)
```

```
[ ]:  order_id  quantity                item_name \
0         1         1      Chips and Fresh Tomato Salsa
1         1         1                      Izze
2         1         1          Nantucket Nectar
3         1         1  Chips and Tomatillo-Green Chili Salsa
4         2         2          Chicken Bowl
5         3         1          Chicken Bowl
6         3         1        Side of Chips
7         4         1        Steak Burrito
8         4         1      Steak Soft Tacos
9         5         1        Steak Burrito

      choice_description  item_price  total_price
0                  NaN          2.39         11.56
1          [Clementine]          3.39         11.56
2              [Apple]          3.39         11.56
3                  NaN          2.39         11.56
4  [Tomatillo-Red Chili Salsa (Hot), [Black Beans...      16.98         16.98
5  [Fresh Tomato Salsa (Mild), [Rice, Cheese, Sou...      10.98         12.67
6                  NaN          1.69         12.67
7  [Tomatillo Red Chili Salsa, [Fajita Vegetables...      11.75         21.00
8  [Tomatillo Green Chili Salsa, [Pinto Beans, Ch...          9.25         21.00
9  [Fresh Tomato Salsa, [Rice, Black Beans, Pinto...          9.25         13.70
```

As you can see, the total price of each order is now listed on every single line.

That makes it easy to calculate the percentage of the total order price that each line represents:

```
[ ]: orders['percent_of_total'] = orders.item_price / orders.total_price
orders.head(10)
```

```
[ ]:  order_id  quantity                item_name \
0         1         1      Chips and Fresh Tomato Salsa
1         1         1                      Izze
2         1         1          Nantucket Nectar
3         1         1  Chips and Tomatillo-Green Chili Salsa
```

4	2	2	Chicken Bowl
5	3	1	Chicken Bowl
6	3	1	Side of Chips
7	4	1	Steak Burrito
8	4	1	Steak Soft Tacos
9	5	1	Steak Burrito

	choice_description	item_price	total_price \
0	NaN	2.39	11.56
1	[Clementine]	3.39	11.56
2	[Apple]	3.39	11.56
3	NaN	2.39	11.56
4	[Tomatillo-Red Chili Salsa (Hot), [Black Beans...	16.98	16.98
5	[Fresh Tomato Salsa (Mild), [Rice, Cheese, Sou...	10.98	12.67
6	NaN	1.69	12.67
7	[Tomatillo Red Chili Salsa, [Fajita Vegetables...	11.75	21.00
8	[Tomatillo Green Chili Salsa, [Pinto Beans, Ch...	9.25	21.00
9	[Fresh Tomato Salsa, [Rice, Black Beans, Pinto...	9.25	13.70

	percent_of_total
0	0.206747
1	0.293253
2	0.293253
3	0.206747
4	1.000000
5	0.866614
6	0.133386
7	0.559524
8	0.440476
9	0.675182

## 1.22 20. Select a slice of rows and columns

Let's take a look at another dataset:

```
[ ]: titanic.head()
```

```
[ ]: PassengerId  Survived  Pclass \
0            1         0         3
1            2         1         1
```

2	3	1	3
3	4	1	1
4	5	0	3

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

This is the famous Titanic dataset, which shows information about passengers on the Titanic and whether or not they survived.

If you wanted a numerical summary of the dataset, you would use the `describe()` method:

```
[ ]: titanic.describe()
```

```
[ ]:
      PassengerId   Survived  Pclass     Age     SibSp  \
count    891.000000   891.000000   891.000000  714.000000  891.000000
mean       446.000000     0.383838     2.308642   29.699118    0.523008
std       257.353842     0.486592     0.836071   14.526497    1.102743
min         1.000000     0.000000     1.000000     0.420000    0.000000
25%       223.500000     0.000000     2.000000   20.125000    0.000000
50%       446.000000     0.000000     3.000000   28.000000    0.000000
75%       668.500000     1.000000     3.000000   38.000000    1.000000
max       891.000000     1.000000     3.000000   80.000000    8.000000
```

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400

```

50%      0.000000    14.454200
75%      0.000000    31.000000
max       6.000000   512.329200

```

However, the resulting DataFrame might be displaying more information than you need.

If you wanted to filter it to only show the “five-number summary”, you can use the `loc` accessor and pass it a slice of the “min” through the “max” row labels:

```
[ ]: titanic.describe().loc['min':'max']
```

```
[ ]:
      PassengerId  Survived  Pclass     Age  SibSp  Parch    Fare
min            1.0        0.0      1.0   0.420    0.0    0.0    0.0000
25%           223.5        0.0      2.0  20.125    0.0    0.0    7.9104
50%           446.0        0.0      3.0  28.000    0.0    0.0   14.4542
75%           668.5        1.0      3.0  38.000    1.0    0.0   31.0000
max           891.0        1.0      3.0  80.000    8.0    6.0  512.3292

```

And if you’re not interested in all of the columns, you can also pass it a slice of column labels:

```
[ ]: titanic.describe().loc['min':'max', 'Pclass':'Parch']
```

```
[ ]:
      Pclass     Age  SibSp  Parch
min      1.0   0.420    0.0    0.0
25%      2.0  20.125    0.0    0.0
50%      3.0  28.000    0.0    0.0
75%      3.0  38.000    1.0    0.0
max      3.0  80.000    8.0    6.0

```

## 1.23 21. Reshape a MultiIndexed Series

The Titanic dataset has a “Survived” column made up of ones and zeros, so you can calculate the overall survival rate by taking a mean of that column:

```
[ ]: titanic.Survived.mean()
```

```
[ ]: 0.3838383838383838
```

If you wanted to calculate the survival rate by a single category such as “Sex”, you would use a `groupby()`:

```
[ ]: titanic.groupby('Sex').Survived.mean()
```

```
[ ]: Sex
      female    0.742038
      male      0.188908
      Name: Survived, dtype: float64
```

And if you wanted to calculate the survival rate across two different categories at once, you would `groupby()` both of those categories:

```
[ ]: titanic.groupby(['Sex', 'Pclass']).Survived.mean()
```

```
[ ]: Sex      Pclass
      female  1          0.968085
           2          0.921053
           3          0.500000
      male    1          0.368852
           2          0.157407
           3          0.135447
      Name: Survived, dtype: float64
```

This shows the survival rate for every combination of Sex and Passenger Class. It's stored as a `MultIndexed Series`, meaning that it has multiple index levels to the left of the actual data.

It can be hard to read and interact with data in this format, so it's often more convenient to reshape a `MultIndexed Series` into a `DataFrame` by using the `unstack()` method:

```
[ ]: titanic.groupby(['Sex', 'Pclass']).Survived.mean().unstack()
```

```
[ ]: Pclass      1          2          3
      Sex
      female  0.968085  0.921053  0.500000
      male    0.368852  0.157407  0.135447
```

This `DataFrame` contains the same exact data as the `MultIndexed Series`, except that now you can interact with it using familiar `DataFrame` methods.

## 1.24 22. Create a pivot table

If you often create `DataFrames` like the one above, you might find it more convenient to use the `pivot_table()` method instead:

```
[ ]: titanic.pivot_table(index='Sex', columns='Pclass', values='Survived',
      ↪aggfunc='mean')
```



```
[ ]: Pclass      1      2      3
      Sex
      female  0.968085  0.921053  0.500000
      male    0.368852  0.157407  0.135447
```

With a pivot table, you directly specify the index, the columns, the values, and the aggregation function.

An added benefit of a pivot table is that you can easily add row and column totals by setting `margins=True`:

```
[ ]: titanic.pivot_table(index='Sex', columns='Pclass', values='Survived',
    ↪aggfunc='mean',
    margins=True)
```

```
[ ]: Pclass      1      2      3      All
      Sex
      female  0.968085  0.921053  0.500000  0.742038
      male    0.368852  0.157407  0.135447  0.188908
      All     0.629630  0.472826  0.242363  0.383838
```

This shows the overall survival rate as well as the survival rate by Sex and Passenger Class.

Finally, you can create a cross-tabulation just by changing the aggregation function from “mean” to “count” :

```
[ ]: titanic.pivot_table(index='Sex', columns='Pclass', values='Survived',
    ↪aggfunc='count',
    margins=True)
```

```
[ ]: Pclass      1      2      3      All
      Sex
      female    94     76    144    314
      male     122    108    347    577
      All      216    184    491    891
```

This shows the number of records that appear in each combination of categories.

## 1.25 23. Convert continuous data into categorical data

Let’ s take a look at the Age column from the Titanic dataset:

```
[ ]: titanic.Age.head(10)
```

```
[ ]: 0    22.0
      1    38.0
      2    26.0
      3    35.0
      4    35.0
      5     NaN
      6    54.0
      7     2.0
      8    27.0
      9    14.0
      Name: Age, dtype: float64
```

It's currently continuous data, but what if you wanted to convert it into categorical data?

One solution would be to label the age ranges, such as “child”, “young adult”, and “adult”. The best way to do this is by using the `cut()` function:

```
[ ]: pd.cut(titanic.Age, bins=[0, 18, 25, 99], labels=['child', 'young adult', 'adult']).head(10)
```

```
[ ]: 0    young adult
      1         adult
      2         adult
      3         adult
      4         adult
      5          NaN
      6         adult
      7         child
      8         adult
      9         child
      Name: Age, dtype: category
      Categories (3, object): [child < young adult < adult]
```

This assigned each value to a bin with a label. Ages 0 to 18 were assigned the label “child”, ages 18 to 25 were assigned the label “young adult”, and ages 25 to 99 were assigned the label “adult”.

Notice that the data type is now “category”, and the categories are automatically ordered.

## 1.26 24. Change display options

Let's take another look at the Titanic dataset:

```
[ ]: titanic.head()
```

```
[ ]: PassengerId  Survived  Pclass  \
0             1         0        3
1             2         1        1
2             3         1        3
3             4         1        1
4             5         0        3
```

```
                                Name    Sex  Age  SibSp  \
0                Braund, Mr. Owen Harris   male  22.0    1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0    1
2                Heikkinen, Miss. Laina   female  26.0    0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)  female  35.0    1
4                Allen, Mr. William Henry   male  35.0    0
```

```
    Parch    Ticket   Fare Cabin Embarked
0      0  A/5 21171   7.2500   NaN        S
1      0  PC 17599  71.2833   C85        C
2      0 STON/O2. 3101282   7.9250   NaN        S
3      0    113803  53.1000  C123        S
4      0    373450   8.0500   NaN        S
```

Notice that the Age column has 1 decimal place and the Fare column has 4 decimal places. What if you wanted to standardize the display to use 2 decimal places?

You can use the `set_option()` function:

```
[ ]: pd.set_option('display.float_format', '{:.2f}'.format)
```

The first argument is the name of the option, and the second argument is a Python format string.

```
[ ]: titanic.head()
```

```
[ ]: PassengerId  Survived  Pclass  \
0             1         0        3
1             2         1        1
2             3         1        3
3             4         1        1
4             5         0        3
```

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.00	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.00	1	
2	Heikkinen, Miss. Laina	female	26.00	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.00	1	
4	Allen, Mr. William Henry	male	35.00	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.25	NaN	S
1	0	PC 17599	71.28	C85	C
2	0	STON/O2. 3101282	7.92	NaN	S
3	0	113803	53.10	C123	S
4	0	373450	8.05	NaN	S

You can see that Age and Fare are now using 2 decimal places. Note that this did not change the underlying data, only the display of the data.

You can also reset any option back to its default:

```
[ ]: pd.reset_option('display.float_format')
```

There are many more options you can specify in a similar way.

## 1.27 25. Style a DataFrame

The previous trick is useful if you want to change the display of your entire notebook. However, a more flexible and powerful approach is to define the style of a particular DataFrame.

Let's return to the stocks DataFrame:

```
[ ]: stocks
```

```
[ ]:
      Date    Close  Volume Symbol
0 2016-10-03   31.50  14070500   CSCD
1 2016-10-03  112.52  21701800   AAPL
2 2016-10-03   57.42  19189500   MSFT
3 2016-10-04  113.00  29736800   AAPL
4 2016-10-04   57.24  20085900   MSFT
5 2016-10-04   31.35  18460400   CSCD
6 2016-10-05   57.64  16726400   MSFT
7 2016-10-05   31.59  11808600   CSCD
8 2016-10-05  113.05  21453100   AAPL
```

We can create a dictionary of format strings that specifies how each column should be formatted:

```
[ ]: format_dict = {'Date': '{:%m/%d/%y}', 'Close': '${:.2f}', 'Volume': '{:,}'}

```

And then we can pass it to the DataFrame's `style.format()` method:

```
[ ]: stocks.style.format(format_dict)

```

```
[ ]: <pandas.io.formats.style.Styler at 0x7fbd42ce9b38>

```

Notice that the Date is now in month-day-year format, the closing price has a dollar sign, and the Volume has commas.

We can apply more styling by chaining additional methods:

```
[ ]: (stocks.style.format(format_dict)
      .hide_index()
      .highlight_min('Close', color='red')
      .highlight_max('Close', color='lightgreen')
      )

```

```
[ ]: <pandas.io.formats.style.Styler at 0x7fbd1074ad30>

```

We've now hidden the index, highlighted the minimum Close value in red, and highlighted the maximum Close value in green.

Here's another example of DataFrame styling:

```
[ ]: (stocks.style.format(format_dict)
      .hide_index()
      .background_gradient(subset='Volume', cmap='Blues')
      )

```

```
[ ]: <pandas.io.formats.style.Styler at 0x7fbd430b7d30>

```

The Volume column now has a background gradient to help you easily identify high and low values.

And here's one final example:

```
[ ]: (stocks.style.format(format_dict)
      .hide_index()
      .bar('Volume', color='lightblue', align='zero')
      .set_caption('Stock Prices from October 2016')
      )

```

```
[ ]: <pandas.io.formats.style.Styler at 0x7fbd10772240>
```

There's now a bar chart within the Volume column and a caption above the DataFrame. Note that there are many more options for how you can style your DataFrame.

## 1.28 Bonus: Profile a DataFrame

Let's say that you've got a new dataset, and you want to quickly explore it without too much work. There's a separate package called [pandas-profiling](#) that is designed for this purpose.

First you have to install it using conda or pip. Once that's done, you import `pandas_profiling`:

```
[ ]: import pandas_profiling
```

Then, simply run the `ProfileReport()` function and pass it any DataFrame. It returns an interactive HTML report:

- The first section is an overview of the dataset and a list of possible issues with the data.
- The next section gives a summary of each column. You can click “toggle details” for even more information.
- The third section shows a heatmap of the correlation between columns.
- And the fourth section shows the head of the dataset.

```
[ ]: pandas_profiling.ProfileReport(titanic)
```

```
[ ]: <pandas_profiling.ProfileReport at 0x7fbd3065d630>
```

1.28.1 Want more tricks? Watch [21 more pandas tricks](#) or [Read the notebook](#)

© 2019 [Data School](#). All rights reserved.